

SANJITH KUMAR

sanjithkumar.in ◇ [GitHub](#)

(+91) 9113944058 ◇ Bengaluru, IN

sanjithkumar048@gmail.com ◇ [LinkedIn](#)

SUMMARY

3 years across two companies: data pipelines and full stack modernization at Carelon, then agentic AI platforms at Deloitte. At Carelon, built PySpark/Kafka eligibility feeds and rebuilt a legacy enrollment app in Angular + Spring Boot. At Deloitte, shipped three enterprise platforms – LangGraph orchestrators, RAG pipelines, FastAPI backends, React frontends, and a custom evals suite across **18 behavior probes**. One PoC closed **\$15M**. *Full stack end-to-end – no hand-offs*.

SKILLS

AI / LLM	LangGraph (StateGraph, ToolNode), LangChain, OpenAI API (GPT-4.5, GPT-5), Gemini, Anthropic Claude, RAG (pgvector, ChromaDB, hybrid BM25 + dense), Prompt Engineering, Multi-Agent Systems, MCP, ReAct, Langfuse
Languages	Python, JavaScript, TypeScript, Java, SQL
Data Engineering	PySpark, Apache Kafka, AWS Glue, ETL Orchestration, Polars
Back-End	FastAPI, Node.js (Express), Django, Spring Boot
Databases	PostgreSQL (pgvector), Oracle SQL, MySQL, SQLAlchemy, Prisma, Drizzle ORM
Front-End	React.js, Next.js, Angular, TypeScript, Vite, Tailwind CSS, Zustand
Cloud & DevOps	AWS (EKS/KEDA, EC2, S3, Lambda, Glue), Docker, Kubernetes, CI/CD
Testing	pytest, Playwright, integration testing

EXPERIENCE

Associate Engineer

Deloitte USI

April 2025 – Present

Bengaluru, IN

- **Member One – Agentic Enrollment Orchestrator:** Built a LangGraph orchestrator binding **16 tools** across file management, configuration, and planning in a cyclic StateGraph loop – weighted planning engine (priority, line-of-business, size, age), capacity-aware execution scheduler, and custom tool-call logging for traceability. Scaled EKS pods via KEDA tied to queue depth rather than CPU, cutting idle overhead on burst loads.
- **Member One – Evals & Prompt Quality:** Lifted triage accuracy **25%** across 5 prompt revisions using a custom two-tier evals suite: deterministic unit tests via fake stubs (no live service dependency) and **18 integration behavior probes** scoring tool selection, guardrail compliance, output hygiene, and task success. Per-case stability scoring surfaced prompts flipping between pass and fail across runs – used as the *prompt-quality meter* throughout development.
- **Polars Ingestion Core:** Parsed **100–400 enrollment files daily** across EDI 834, WCF, JSON, and XML formats, with a family-grouping algorithm that derives subscriber keys per format to keep dependents correctly linked across households.
- **Healthcare Enrollment Platform (DEP):** Built the AI layer (business-rule extraction, semantic guardrails, agentic validation), FastAPI backend, and React.js frontend for a healthcare cloud-migration PoC that drove a **\$15M client commitment**. Covered the backend with pytest (unit + integration) and validated enrollment flows end-to-end with Playwright.
- **Enterprise Estimation Platform (DAF):** Built a LangGraph + OpenAI multi-agent estimation pipeline, FastAPI backend, and React.js frontend. Refactored a flat JSON store into a **normalized 5-dimensional PostgreSQL schema** with RBAC. Built the retrieval layer on pgvector with **hybrid search** (dense OpenAI embeddings + BM25), benchmarking latency and recall against ChromaDB before migrating. Integrated Langfuse to trace per-agent spans, tool-call latency, and token costs; iterated prompts across GPT-4.5 and GPT-5 against fixed test cases to improve estimation quality.

Associate Software Engineer

Carelon Global Solutions

July 2023 – March 2025

Bengaluru, IN

- **Member Outbound Datalayer (MOD):** Built an event-driven pipeline (PySpark, AWS Glue, Kafka) for high-volume real-time eligibility feeds – cut Mainframe processing time by **40%** and lifted data accuracy by **35%**.
- **Full-Stack Modernization:** Rebuilt the enrollment app in Angular + Spring Boot; rewrote slow queries with targeted indexes (**25% faster**) and replaced bulk loads with paginated lazy views. Added pytest coverage across backend services.
- **ETL Orchestration:** Cut vendor onboarding from 3–4 weeks to 3–4 days (**85% faster**) with a config-driven PySpark/Glue/PostgreSQL ETL layer, replacing manual file transfers with a self-service model.

PROJECTS

- **Agentic Financial Advisor:** Stateful multi-agent assistant (Python, LangGraph, LangChain) – cyclic StateGraph hard-capped at 3 review loops to prevent runaway cycles. ReAct Researcher pulls live market data via Google Search API; Reviewer audits and self-corrects before rendering to Streamlit. ([Repo Link](#))

EDUCATION

Bachelor of Engineering in Computer Science

Cambridge Institute of Technology

2019 – 2023

7.56 CGPA

Relevant Coursework: Full-Stack Development, Database Management, Object-Oriented Programming, Machine Learning