

SANJITH KUMAR

sanjithkumar.in ◇ [GitHub](#)

(+91) 9113944058 ◇ Bengaluru, IN

sanjithkumar048@gmail.com ◇ [LinkedIn](#)

SUMMARY

3 years across two companies: data pipelines and full stack modernization at Carelon, then agentic AI platforms at Deloitte. At Carelon, built PySpark/Kafka eligibility feeds and rebuilt a legacy enrollment app in Angular + Spring Boot. At Deloitte, shipped three enterprise platforms – LangGraph orchestrators, RAG pipelines, FastAPI backends, React frontends, and a custom evals suite across **18 behavior probes**. One PoC closed **\$15M**. *Full stack end-to-end – no hand-offs.*

SKILLS

AI / LLM	LangGraph (StateGraph, ToolNode), LangChain, OpenAI API (GPT-4.5, GPT-5), Gemini, Anthropic Claude, RAG (pgvector, ChromaDB, hybrid BM25 + dense), Prompt Engineering, Multi-Agent Systems, MCP, ReAct, Langfuse
Languages	Python, JavaScript, TypeScript, Java, SQL
Data Engineering	PySpark, Apache Kafka, AWS Glue, ETL Orchestration, Polars
Back-End	FastAPI, Node.js (Express), Django, Spring Boot
Databases	PostgreSQL (pgvector), Oracle SQL, MySQL, SQLAlchemy, Prisma, Drizzle ORM
Front-End	React.js, Next.js, Angular, TypeScript, Vite, Tailwind CSS, Zustand
Cloud & DevOps	AWS (ECS, EKS/KEDA, EC2, Fargate, S3, Lambda, Glue), Docker, Kubernetes, CI/CD
Testing	pytest, Playwright, integration testing
AI Dev Tools	Claude Code, GitHub Copilot, Devin

EXPERIENCE

Associate Engineer

Deloitte USI

April 2025 – Present

Bengaluru, IN

- **Member One – Agentic Enrollment Orchestrator:** Built a LangGraph orchestrator binding **16 tools** in a cyclic StateGraph loop – raw member data into **MongoDB** for intake, promoted to **PostgreSQL** post-enrichment; weighted planning engine (priority, line-of-business, size, age), scheduler with per-job capacity checks, and custom tool-call logging. Scaled EKS pods via KEDA on queue depth, not CPU – idle pods stop burning compute between bursts.
- **Member One – Evals & Prompt Quality:** Lifted triage accuracy **25%** across 5 prompt revisions using a custom two-tier evals suite: deterministic unit tests via fake stubs (no live service dependency) and **18 integration behavior probes** scoring tool selection, guardrail compliance, output hygiene, and task success. Per-case stability scoring surfaced prompts flipping between pass and fail across runs – used as the *prompt-quality meter* throughout development.
- **Claude Code – Dev Toolchain:** Used Claude Code for test scaffolding, debugging, a **code-review agent** enforcing **SOLID** and **YAGNI** principles, and a **security audit agent** that scanned the codebase for vulnerabilities and generated reports.
- **Polars Ingestion Core:** Parsed **100–400 enrollment files daily** across **EDI 834, WCF, JSON, and XML**. Each format needed its own subscriber-key derivation logic – without it, dependents end up detached from the wrong subscriber or duplicated across households. Built a custom allocator skill that reads file priority and size at dispatch and sets vCPU and memory per **AWS Batch** job running on a **Fargate** compute environment; low-priority small files stop over-provisioning, large burst files stop OOM-ing mid-job.
- **Healthcare Enrollment Platform (DEP):** Built the AI layer (business-rule extraction, semantic guardrails, agentic validation), FastAPI backend, and React.js frontend for a healthcare cloud-migration PoC that closed a **\$15M deal**.
- **Enterprise Estimation Platform (DAF):** Built a LangGraph + OpenAI multi-agent estimation pipeline, FastAPI backend, and React.js frontend. Refactored a flat JSON store into a **normalized 5-dimensional PostgreSQL schema** with RBAC. Built the retrieval layer on pgvector with **hybrid search** (dense OpenAI embeddings + BM25), benchmarking latency and recall against ChromaDB before migrating. Integrated Langfuse to trace per-agent spans, tool-call latency, and token costs; ran prompt iterations across GPT-4.5 and GPT-5 against fixed test cases.

Associate Software Engineer

Carelon Global Solutions

July 2023 – March 2025

Bengaluru, IN

- **Member Outbound Datalayer (MOD):** Built an event-driven pipeline (PySpark, AWS Glue, Kafka) processing real-time eligibility feeds – **3M records daily**, peaking at **20M** at month-end. Schema validation caught malformed feeds before they hit downstream systems. Cut Mainframe processing time by **40%** and improved data accuracy by **35%**.
- **Full-Stack Modernization:** Rebuilt the enrollment app in Angular + Spring Boot; rewrote slow queries with targeted indexes (**25% faster**) and replaced bulk loads with paginated lazy views. Added pytest coverage across backend services.
- **ETL Orchestration:** Cut vendor onboarding from 3–4 weeks to 3–4 days (**85% faster**) across **15–20 active vendors**, with new ones added monthly. A DB-driven config layer replaced manual file transfers – onboarding a vendor meant updating a config record, not writing pipeline code. Worked directly with vendors to map their source formats into the ingestion schema.

PROJECTS

- **Agentic Financial Advisor:** Stateful multi-agent assistant (Python, LangGraph, LangChain) – cyclic StateGraph hard-capped at 3 review loops to prevent runaway cycles. ReAct Researcher pulls live market data via Google Search API; Reviewer audits and self-corrects before rendering to Streamlit. ([Repo Link](#))

EDUCATION

Bachelor of Engineering in Computer Science

Cambridge Institute of Technology

2019 – 2023

7.56 CGPA

Relevant Coursework: Full-Stack Development, Database Management, Object-Oriented Programming, Machine Learning